

# Machine learning: a hands-on introduction

Filippo Biscarini (CNR, Milan, Italy)

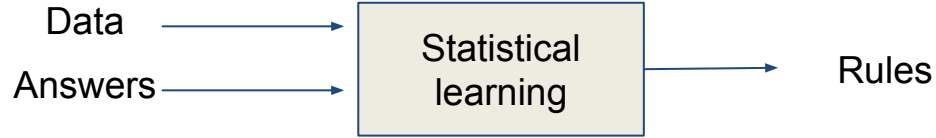
[filippo.biscarini@cnr.it](mailto:filippo.biscarini@cnr.it)



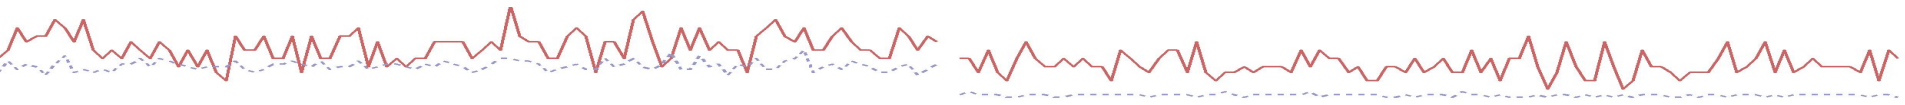
# Supervised learning




# What is learning?




- building a statistical model for **predicting** an **output** based on one or more **inputs**
- statistical learning model is **trained** rather than explicitly programmed



# What is learning?



1. Input data (e.g. genome variants, metabolites)
2. Output examples (e.g. disease status, biological characteristics)
3. Performance measure: how well is the algorithm working → adjustment steps  
→ **learning**



# You can do (statistical) learning in your head!



- The current price of an electric Tesla is \$100,000
- The price of an electric Tesla next year will be \$90,000
- The price of an electric Tesla in two years will be \$81,000
- The price of an electric Tesla in three years will be \$72,900
- How much will an electric Tesla cost in five years?



# You can do (statistical) learning in your head!

## TRAINING DATA

- The current price of an electric Tesla is \$100,000
- The price of an electric Tesla next year will be \$90,000
- The price of an electric Tesla in two years will be \$81,000
- The price of an electric Tesla in three years will be \$72900
- How much will an electric Tesla cost in five years?

## NEW, UNKNOWN DATA

$\text{PRICE} = \text{PRICE}_0 \cdot (1 - 0.10)^{\text{YEARS}}$

PRICE IN FIVE YEARS = \$59,049

MATHEMATICAL  
MODEL

PREDICTION




# The truth about learning

- Yesterday: **n. of cases** =  $2^{(\text{days}-1)}$
- Today: **price** =  $\text{price}_0 * (1-0.10)^{\text{years}}$

Is this what ML is learning?



# The truth about learning



- (known) **quadratic function** (blue line)
- approximated with **machine learning** (orange line)

**!! ML is not learning  $y = x^2$  !!**



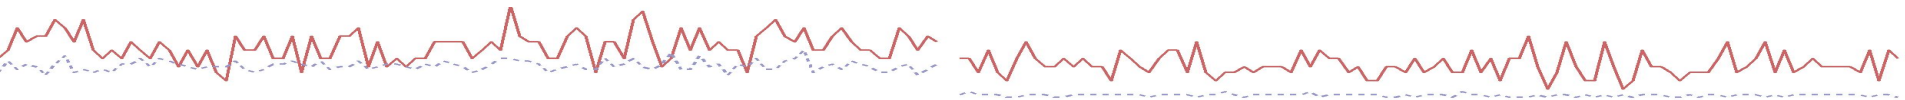
# Why supervised?

## Training examples

measured variables / features  
on  $n$  examples

$$\begin{bmatrix} 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix} = \begin{bmatrix} 0.12 & 1.5 & \dots & 0.9 \\ 2.05 & 0.95 & \dots & 1.1 \\ \vdots & \vdots & \ddots & \vdots \\ 3.5 & 0.88 & \dots & 1.75 \end{bmatrix}$$

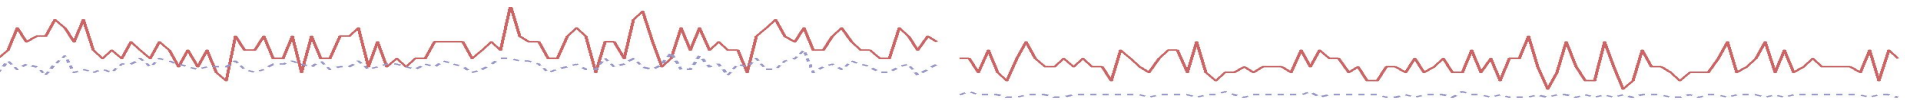
labels / target  
variables (e.g.  
phenotypes) on  $n$   
examples



# Unsupervised learning

measured variables / features  
on  $n$  examples

$$\cancel{\begin{bmatrix} 1 \\ 0 \\ \vdots \\ 1 \end{bmatrix}} = \begin{bmatrix} 0.12 & 1.5 & \dots & 0.9 \\ 2.05 & 0.95 & \dots & 1.1 \\ \vdots & \vdots & \ddots & \vdots \\ 3.5 & 0.88 & \dots & 1.75 \end{bmatrix}$$



# The steps of a supervised learning problem



1. Collect the data
2. EDA and data preparation
3. Training a model on the data
4. Evaluate model performance
5. Improve model performance

80% of the time!

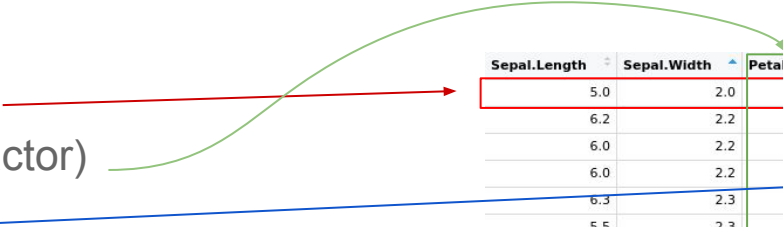
rinse and repeat!

model deployment



# A little ML jargon

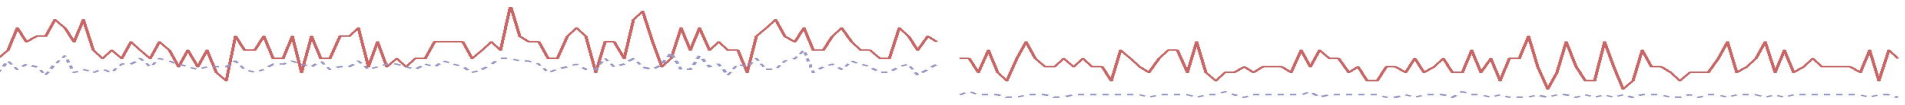
- **example** (record, observation)
- **feature** (independent variable, factor)
- **label** (dependent variable)
- **method**: the statistical method used for a problem
- **model**: the modelling of the problem (e.g. which features to include and how)
- **algorithm**: the technique by which the method is applied to the model and solved
- **training data**: data on which the ML algorithm is trained



| Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species    |
|--------------|-------------|--------------|-------------|------------|
| 5.0          | 2.0         | 3.5          | 1.0         | versicolor |
| 6.2          | 2.2         | 4.5          | 1.5         | versicolor |
| 6.0          | 2.2         | 4.0          | 1.0         | versicolor |
| 6.0          | 2.2         | 5.0          | 1.5         | virginica  |
| 6.3          | 2.3         | 4.4          | 1.3         | versicolor |
| 5.5          | 2.3         | 4.0          | 1.3         | versicolor |
| 5.0          | 2.3         | 3.3          | 1.0         | versicolor |
| 4.5          | 2.3         | 1.3          | 0.3         | setosa     |
| 5.5          | 2.4         | 3.8          | 1.1         | versicolor |
| 5.5          | 2.4         | 3.7          | 1.0         | versicolor |
| 4.9          | 2.4         | 3.3          | 1.0         | versicolor |
| 6.7          | 2.5         | 5.8          | 1.8         | virginica  |
| 6.3          | 2.5         | 4.9          | 1.5         | versicolor |

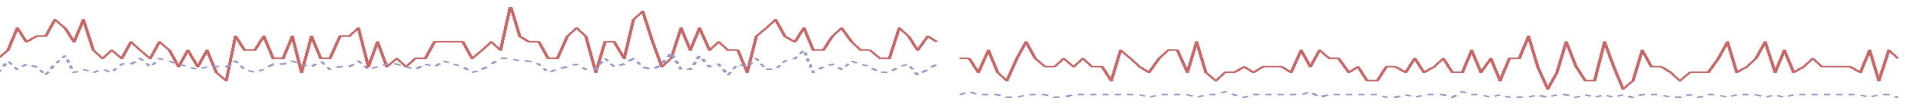


# Regression and classification



# Supervised learning problems

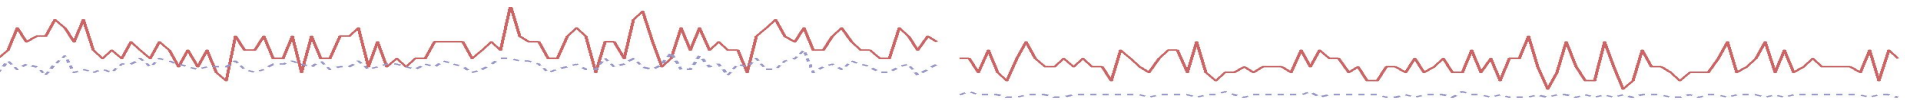
- Regression problems
- Classification problems



# Supervised learning problems

- Regression (**predictive**) problems
  - target (continuous) variable, output
- Classification (**predictive**) problems
  - label, class (qualitative variable): binomial, multinomial, ordinal, nominal

*“given a set of data, the learning algorithm attempts to optimize a function (the model) to find the combination of feature values that result in the target output”*



# Supervised learning problems

- Regression (**predictive**) problems
  - target (continuous) variable, output
- Classification (**predictive**) problems
  - label, class (qualitative variable): binomial, multinomial, ordinal, nominal


## Predict:

- the future (forecasting, prognosis)
- the unknown/unseen (e.g. sick/healthy, genetic predisposition etc.)
- real time (e.g. control traffic lights at rush hours)
- the past (e.g. when something happened, like conception date based on hormone levels)



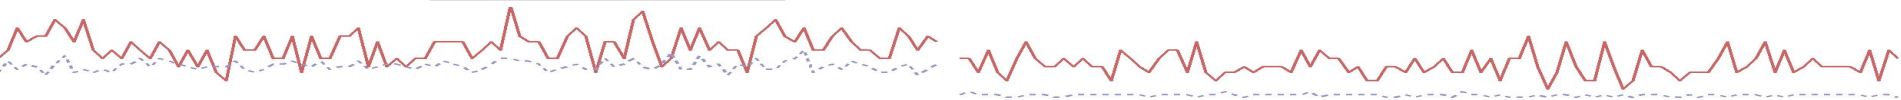


# Supervised learning problems

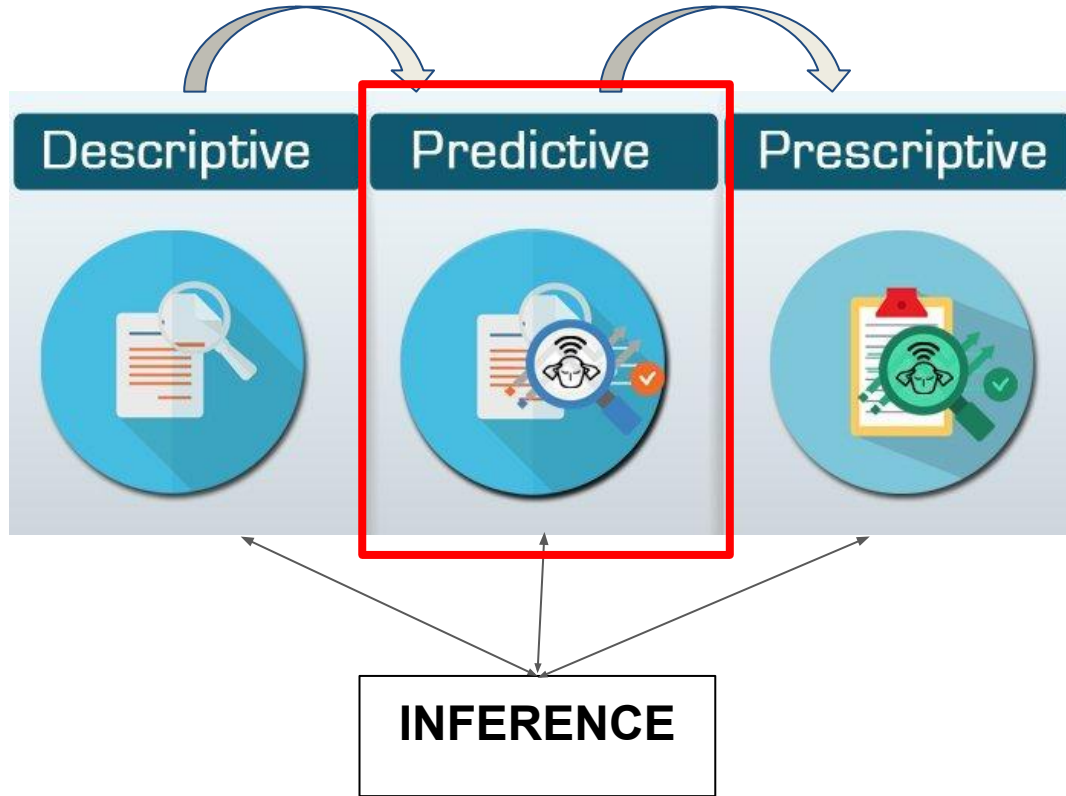


- Know the past
- Predict the future
- Act consequently

**INFERENCE**

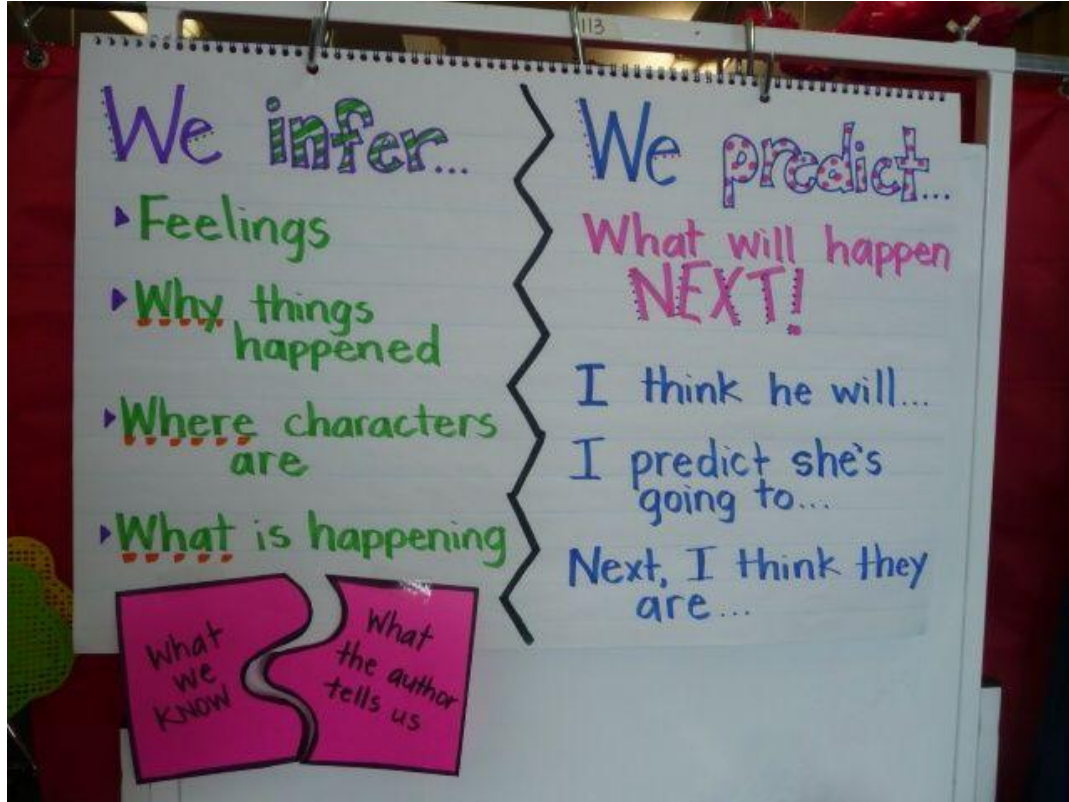


# Supervised learning problems



- Know the past
- Predict the future
- Act consequently


# Inference vs Prediction



- different statistical problems
- different objectives, different rules ... different ballparks
- inference is in general more difficult than prediction

# What causes what?

- Typical inferential problem
- Correlation is not causation!
  - already since Pearson! (early 1900's)
  - [spurious correlations generator](#)
  - *apophenia* (from Greek, "to seem")



# What causes what?


- Typical inferential problem
- Correlation is not causation!

## Reverse causation

- n. of firemen  $\longleftrightarrow$  size of the fire (the more the firemen, the bigger the fire?)
- smoking  $\longleftrightarrow$  depression (smoking causes depression?)

## Missing variable

- ice cream consumption  $\longleftrightarrow$  n. of sunburn cases
- buying lighters  $\longleftrightarrow$  lung cancer



# What causes what?


- Typical inferential problem
- Correlation is not causation!

## Reverse causation


- n. of firemen  $\longleftrightarrow$  size of the fire (the more the firemen, the bigger the fire?)
- smoking  $\longleftrightarrow$  depression (smoking causes depression?)

## Missing variable

- ice cream consumption  $\longleftrightarrow$  n. of sunburn cases
- buying lighters  $\longleftrightarrow$  lung cancer




**Q: reverse causation or missing variable?**






# What causes what?

- Also regression!
  - **Galton**: predict the height of children from the height of their parents
  - But also the height of parents could be predicted from that of their children! (??)→no possible cause-effect relationship!



[Han, Ma, Zhu 2015]

# What causes what?




spurious regression: can you spot the mistakes?

MY RESEARCH SUGGESTS THE EXISTENCE OF AN 800-METER SPHERE THAT TASTES OKAY.





# The probabilistic nature of cause (it gets trickier)




- you turn the key and the car engine starts: clear cause-effect link
- you take an ibuprofen pill and your headache goes away: how do you know it wouldn't have gone anyway? → **counterfactual**
- smoking→lung cancer: my uncle always smoked and did not have lung cancer! my aunt never smoked and got lung cancer! →if you smoke it is **more likely** that you get lung cancer

$X \rightarrow Y$ : X increases/decreases the chances that Y happens



# The probabilistic nature of cause (it gets trickier)



**RCT:** randomized  
(clinical) trial

|                      | placebo % | statin % | relative risk reduction |
|----------------------|-----------|----------|-------------------------|
| heart attack         | 11.8      | 8.7      | 27%                     |
| stroke               | 5.7       | 4.3      | 25%                     |
| death from any cause | 14.7      | 12.9     | 13%                     |

[Heart Protection Study, 2002: [https://www.thelancet.com/journals/lancet/article/PIIS0140-6736\(02\)09327-3/fulltext](https://www.thelancet.com/journals/lancet/article/PIIS0140-6736(02)09327-3/fulltext)]

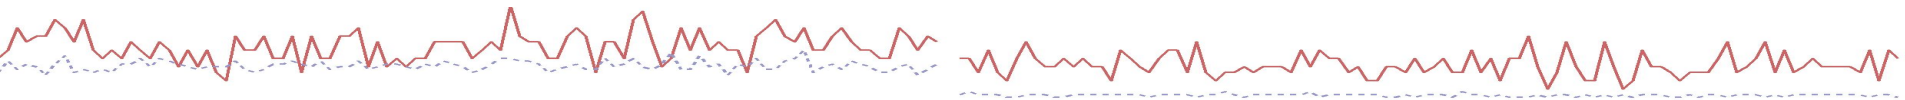


# What causes what?

- **experiments** (RCT, A/B testing etc.)

## without experiments

- post hoc, ergo propter hoc
- strength of correlation
- consistency (different datasets, studies etc.)
- dose-response relationship
- analysis “ceteris paribus” (controlling for confounders, stratifiers etc.)
- Bayesian Networks
- ...



# What causes what? ML perspective

**“observational data alone does not provide causal insight”** (Pearl and Mackenzie 2018\*)

ML can help with causal questions:

- studying associations between variables: starting point for causal hypotheses
- estimating causal effects: quantify the dose-response relationship
- learning causal models: tools to reason about interventions and counterfactuals
- learning causal graphs: direction of causal relationships
- etc.

**The gist of it: if you can predict  $y$  using  $x$ , then there probably is a causal relationship between  $x$  and  $y$**

\*Pearl, J. and Mackenzie, D., 2018. The book of why: the new science of cause and effect. Basic books



# Supervised learning problems

- Regression (**predictive**) problems
- Classification (**predictive**) problems

## Predictive machines!


- Classifiers
- Predictors/Regressors



source:

<https://blog.bigml.com/2013/03/12/machine-learning-from-streaming-data-two-problems-two-solutions-two-concerns-and-two-lessons/>

# Regression


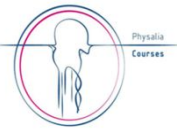


# Regression problems

- the response variable **y** is **quantitative**
- e.g.: *height, weight, yield (milk, crops), blood sugar concentration*
- **y** = **target** (dependent) variable (a.k.a. response, objective variable)
- **X** = matrix of **features** (continuous, categorical)
- **predictor**:  $y = f(x) = \mathbf{P}(\mathbf{X}) \leftarrow$  [predictive machine]



# Regression problems - simple regression



$$\text{happiness} = (\text{intercept}) + \text{beta} * \text{income}$$

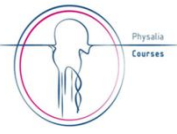
or

$$\text{income} = (\text{intercept}) + \text{beta} * \text{happiness}$$


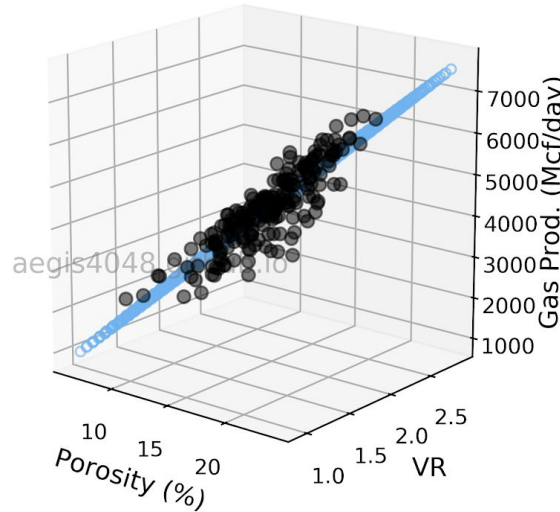
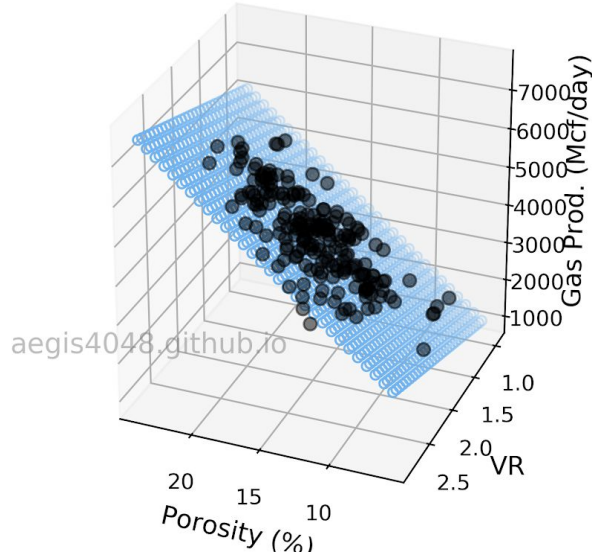
Source: <https://www.scribbr.com/statistics/linear-regression-in-r/>



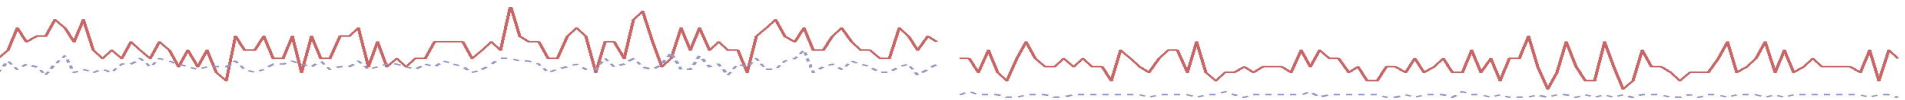
# Regression problems - multiple regression



$$R^2 = 0.79$$



Source: [https://aegis4048.github.io/mutiple\\_linear\\_regression\\_and\\_visualization\\_in\\_python](https://aegis4048.github.io/mutiple_linear_regression_and_visualization_in_python)



# Multiple linear regression

$$y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$$

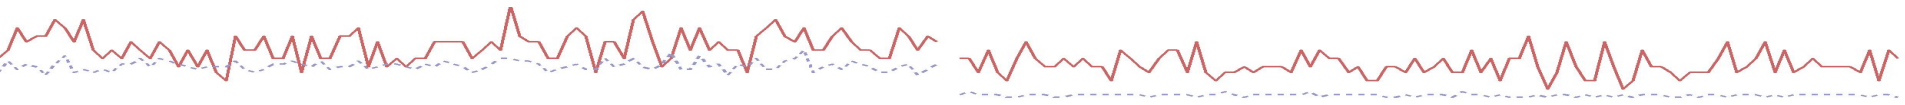
- $y$ : target variable
- $\beta$ 's: model coefficients
- $X$ 's: features (predictors, independent variables, factors)



# Multiple linear regression

$$\mathbf{y} = \beta \mathbf{X} + \mathbf{e}$$

- matrix (compact) notation
- vectors of observations ( $\mathbf{y}$ ), coefficients ( $\beta$ ) and residuals ( $\mathbf{e}$ )
- matrix of features ( $\mathbf{X}$ )



# Multiple linear regression


$$\mathbf{y} = \beta \mathbf{X} + \mathbf{e}$$

estimation of  
coefficients

$$\hat{\mathbf{y}} = \hat{\beta} \mathbf{X}$$

→ predictions!


- matrix (compact) notation
- vectors of observations ( $\mathbf{y}$ ), coefficients ( $\beta$ ) and residuals ( $\mathbf{e}$ )
- matrix of features ( $\mathbf{X}$ )



# Predictions

$$\hat{y} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p$$

with the estimated coefficients  $\beta$  and the feature values  $\mathbf{X}$  we obtain the predicted values  $\hat{y}$




# Predictions

$$\hat{y} = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p$$

with the estimated coefficients  $\beta$  and the feature values  $\mathbf{X}$  we obtain the predicted values  $\hat{y}$

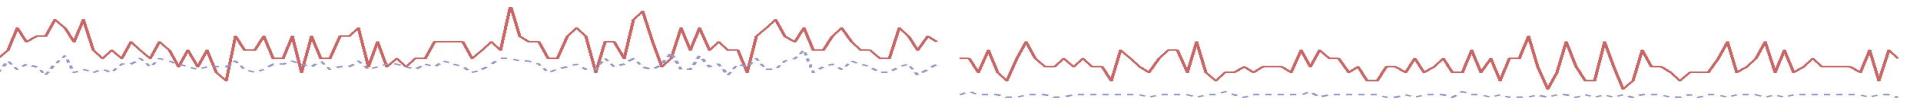
→ **how do we obtain the model coefficients  $\beta$ ?**



# Estimation of model coefficients

$$y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$$


- define a **loss (cost) function**
- **minimise** the loss function



# Estimation of model coefficients

- define a **loss (cost) function**

| observations | predictions                                |
|--------------|--|
| $\mathbf{y}$ | $\hat{\mathbf{y}} = \hat{\beta}\mathbf{X}$ |



difference between observed and  
predicted values






# Estimation of model coefficients

- define a **loss (cost) function**


| observations | predictions                                |
|--------------|--|
| $\mathbf{y}$ | $\hat{\mathbf{y}} = \hat{\beta}\mathbf{X}$ |



difference between observed and  
predicted values



**LEAST SQUARES**



# Estimation of model coefficients


- minimise the **loss (cost) function**

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RSS = \sum_{i=1}^n (y_i - \hat{\beta}_i X_i)^2$$

minimize  $(RSS)$   
 $(\beta)$

**LEAST SQUARES**




# Estimation of model coefficients

- minimise the **loss (cost) function**


$$J(\beta) = \frac{1}{2n} \sum_{i=1}^n (\beta_i X_i - y_i)^2$$

minimize  $J(\beta)$   
 $\beta$

modified  
(normalized) RSS  
function



# Minimise the loss function




Simple linear regression (1 parameter):

$$y = \beta \cdot x$$



# Minimise the loss function



1. loss function

2. model parameters

3. minimum

Simple linear regression (1  
parameter):

$$y = \beta \cdot x$$

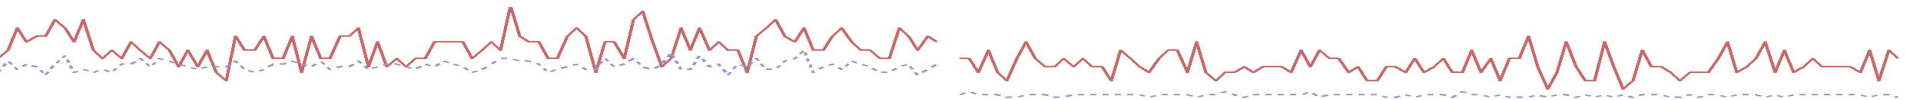




# Minimise the loss function



Multiple linear regression (e.g.

2 parameters):

$$y = \beta_0 + \beta_1 \cdot x$$



# Minimise the loss function




Multiple linear regression (e.g.  
2 parameters):

$$y = \beta_0 + \beta_1 \cdot x$$

Multiple linear regression (> 2  
parameters):

→ m-dimensional hyperspace



# Minimise the loss function

- Demonstration 1.1
- Exercise 1.1

→ 1.introduction\_to\_ml.Rmd





# Minimising the cost function


- the defined cost function is **convex** (it has a minimum!)
- we saw an empirical approach to finding the minimum (manually try some values for the parameters) and the least squares approach

**how do we minimise the cost function in machine learning?**



# Minimising the cost function

- can be minimised by **gradient descent**
- machine learning perspective: gradient descent is a general algorithm to solve models
- alternatively:
  - maximum likelihood
  - (non-)linear least squares



# Loss function: finding the minimum?


Gradient Descent:

minimize  $J(\beta)$   
 $\beta$

1. Start with initial values for  $\beta$  : (initialisation)
2. Change  $\beta$  in the direction of reducing  $J(\beta)$  : (descent)
3. Stop when the minimum is reached : (minimisation)




# Gradient descent




1. starting point (initialisation)
2. find the steepest direction around the starting point
3. take one step in this direction
4. repeat until convergence (local minimum)

Source: Andrew Ng

<https://medium.com/@DBCerigo/on-why-gradient-descent-is-even-needed-25160197a635>



# Gradient descent



Source: Andrew Ng

<https://medium.com/@DBCerigo/on-why-gradient-descent-is-even-needed-25160197a635>

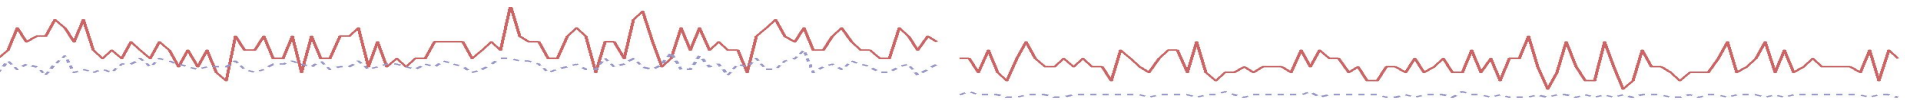
# Gradient descent

$$\beta_j := \beta_j - \alpha \frac{\partial}{\partial \beta_j} J(\beta)$$


assignment  
operator

learning rate (size  
of the steps)


Partial derivative  
of  $J(\beta)$  with  
respect to  $\beta_j$



# Gradient descent



# Gradient descent




$$\beta_j := \beta_j - \alpha \frac{\partial}{\partial \beta_j} J(\beta)$$

- positive slope  $\rightarrow$  reducing the value of  $\beta$  (and the other way around)




# Gradient descent



- $\alpha$  controls the size of the updating step
- small  $\alpha \rightarrow$  slow descent

$$\beta_j := \beta_j - \alpha \frac{\partial}{\partial \beta_j} J(\beta)$$

# Gradient descent




- $\alpha$  controls the size of the updating step
- large  $\alpha \rightarrow$  overshooting: failure to converge

$$\beta_j := \beta_j - \alpha \frac{\partial}{\partial \beta_j} J(\beta)$$



# Gradient descent - recap

- general method to **solve machine learning models** (e.g. multiple linear regression)
- optimise (minimise) the loss function → **optimiser**
- importance of the **learning rate**
- local minimum → **momentum**




# Linear regression - recap

$$1) \quad y = X \cdot \beta + e$$

$$2) \quad \hat{y} = X\hat{\beta}$$

$$3) \quad J(\beta) = \frac{1}{2n} \sum_{i=1}^n \left( y_i - \hat{\beta}_i X_i \right)^2$$



*[minimise  $J(B)$  -  
take derivatives -  
and update  
parameters]*




# Take away message

- **linear regression** from the machine learning perspective → a **predictive machine**
- to be able to make predictions, we first need to **estimate parameter coefficients**
- define a **loss function** and then use **gradient descent** to **minimize it**
- partial derivatives are used to **update** the values of the **parameters**
- **gradient descent** is a general method to minimise the loss (cost) function for a variety of machine learning models



# Measuring performance


- we have our model
- we have estimated the parameters (coefficients) of the model
- we can now get predictions from our predictive machine



# Measuring performance

- we have our model
- we have estimated the parameters (coefficients) of the model
- we can now get predictions from our predictive machine

→ how well are we doing?




# (root) Mean squared error (MSE)

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - f(x_i))^2$$

- average squared difference between predictions and observations

$$RMSE = \sqrt{MSE}$$

- on the same scale as the target variable





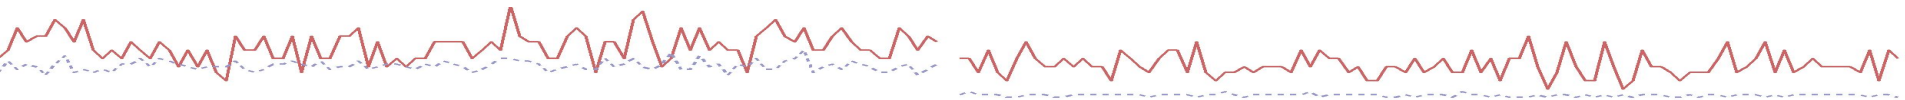
# Mean absolute error (MAE)

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - f(x_i)|$$

- less sensitive to outliers

and the normalized version:

$$NMAE = \frac{MAE}{\bar{y}}$$



# Correlations

- **Pearson's** linear correlation coefficient:


$$\rho_{y,\hat{y}}$$

- **Spearman's** rank correlation coefficient:

$$\rho_{r_y, r_{\hat{y}}}$$



rank variables!



# Measuring performance

- Demonstration 1.2
- Exercise 1.2

→ `introduction_to_ml.Rmd`

